

Sentiment Analysis for Marketing


PHASE -4

DATE	24 October 2023
TEAM ID	Proj-212173-Team-1
PROJECT NAME	Sentiment analysis for Marketing

```
✓ [13] MODEL = 'bert-base-uncased'  
1s tokenizer = BertTokenizer.from_pretrained(MODEL, do_lower_case=True)
```

Downloading (...)okenizer_config.json: 100%  28.0/28.0 [00:00<00:00, 664B/s]
Downloading (...)solve/main/vocab.txt: 100%  232k/232k [00:00<00:00, 819kB/s]
Downloading (...)main/tokenizer.json: 100%  466k/466k [00:00<00:00, 12.5MB/s]
Downloading (...)lve/main/config.json: 100%  570/570 [00:00<00:00, 19.6kB/s]

```
✓ [14] sns.displot(tokens)  
8s plt.xlim([0, 60]);  
plt.xlabel('Token count');
```





```
from torch.utils.data import Dataset, DataLoader
```

```
# Define a custom dataset, more info on how to build custom dataset can be  
# found at https://pytorch.org/tutorials/beginner/data\_loading\_tutorial.html  
class CustomDataset(Dataset):
```

```
    def __init__(  
        self,  
        tweets,  
        labels,  
        tokenizer,  
        max_length  
    ):  
        self.tweets = tweets  
        self.labels = labels  
        self.tokenizer = tokenizer  
        self.max_length = max_length
```

```
    def __len__(self):  
        return len(self.tweets)
```

```
    def __getitem__(self, idx):  
        tweet = self.tweets[idx]  
        label = self.labels[idx]
```

```
    def tokenize(self, tweet):  
        tokenize = self.tokenizer.encode_plus(  
            tweet,  
            add_special_tokens=True,  
            max_length=self.max_length,  
            return_token_type_ids=False,  
            padding='max_length',  
            return_attention_mask=True,  
            return_tensors='pt'  
        )  
        return {  
            'tweet': tweet,  
            'input_ids': tokenize['input_ids'].flatten(),  
            'attention_mask': tokenize['attention_mask'].flatten(),  
            'targets': torch.tensor(label, dtype=torch.long)}
```

✓
0s

```
MAX_LENGTH = 64
TEST_SIZE = 0.1
VALID_SIZE = 0.5
BATCH_SIZE = 16
NUM_WORKERS = 2

train_sampler, test_sampler = train_test_split(df, test_size=TEST_SIZE, random_state=RANDOM_STATE)
valid_sampler, test_sampler = train_test_split(test_sampler, test_size=VALID_SIZE, random_state=RANDOM_STATE)

train_set = CustomDataset(
    train_sampler['text'].to_numpy(),
    train_sampler['labels'].to_numpy(),
    tokenizer,
    MAX_LENGTH
)
test_set = CustomDataset(
    test_sampler['text'].to_numpy(),
    test_sampler['labels'].to_numpy(),
    tokenizer,
    MAX_LENGTH
)
valid_set = CustomDataset(
    valid_sampler['text'].to_numpy(),
    valid_sampler['labels'].to_numpy(),
    tokenizer,
    MAX_LENGTH
)
```

```
train_loader = DataLoader(train_set, batch_size=BATCH_SIZE, num_workers=NUM_WORKERS)
test_loader = DataLoader(test_set, batch_size=BATCH_SIZE, num_workers=NUM_WORKERS)
valid_loader = DataLoader(valid_set, batch_size=BATCH_SIZE, num_workers=NUM_WORKERS)
```

✓
0s

```
from torch import nn
class AirlineSentimentClassifier(nn.Module):

    def __init__(self, num_labels):
        super(AirlineSentimentClassifier, self).__init__()
        self.bert = BertModel.from_pretrained(MODEL)
        self.dropout = nn.Dropout(p=0.2)
        self.classifier = nn.Linear(self.bert.config.hidden_size, num_labels)

    def forward(self, input_ids, attention_mask):
        outputs = self.bert(
            input_ids=input_ids,
            attention_mask=attention_mask
        )
        pooled_output = outputs[1]
        pooled_output = self.dropout(pooled_output)
        out = self.classifier(pooled_output)
        return out
```

✓ [19] 6s model = AirlineSentimentClassifier(len(labels_map))
print(model)

```
# Move tensors to GPU on CUDA enables devices
if device:
    model.cuda()
```

Downloading model.safetensors: 100%  440M/440M [00:04<00:00, 33.2MB/s]

```
AirlineSentimentClassifier(
  (bert): BertModel(
    (embeddings): BertEmbeddings(
      (word_embeddings): Embedding(30522, 768, padding_idx=0)
      (position_embeddings): Embedding(512, 768)
      (token_type_embeddings): Embedding(2, 768)
      (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
    (encoder): BertEncoder(
      (layer): ModuleList(
        (0-11): 12 x BertLayer(
          (attention): BertAttention(
            (self): BertSelfAttention(
              (query): Linear(in_features=768, out_features=768, bias=True)
              (key): Linear(in_features=768, out_features=768, bias=True)
              (value): Linear(in_features=768, out_features=768, bias=True)
              (dropout): Dropout(p=0.1, inplace=False)
            )
            (output): BertSelfOutput(
              (dense): Linear(in_features=768, out_features=768, bias=True)
              (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
              (dropout): Dropout(p=0.1, inplace=False)
            )
          )
        )
      )
    )
  )
)
```

```

        (intermediate): BertIntermediate(
          (dense): Linear(in_features=768, out_features=3072, bias=True)
          (intermediate_act_fn): GELUActivation()
        )
        (output): BertOutput(
          (dense): Linear(in_features=3072, out_features=768, bias=True)
          (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
          (dropout): Dropout(p=0.1, inplace=False)
        )
      )
    )
    (pooler): BertPooler(
      (dense): Linear(in_features=768, out_features=768, bias=True)
      (activation): Tanh()
    )
  )
  (dropout): Dropout(p=0.2, inplace=False)
  (classifier): Linear(in_features=768, out_features=3, bias=True)
)

```

✓ 1s

```

n_epochs = 10
learning_rate = 2e-5

# Loss function
criterion = nn.CrossEntropyLoss()

# Optimizer
optimizer = AdamW(model.parameters(), lr=learning_rate, correct_bias=False)

# Define scheduler
training_steps = len(train_loader)*n_epochs
scheduler = get_linear_schedule_with_warmup(
    optimizer,
    num_warmup_steps=0,
    num_training_steps=training_steps
)

```

⚠ /usr/local/lib/python3.10/dist-packages/transformers/optimization.py:411: FutureWarning: This implementation of AdamW is deprecated. Please use the implementation in torch.nn.optim.AdamW instead. warnings.warn(

```

# Track changes in validation Loss
valid_loss_min = np.Inf

for epoch in range(1, n_epochs+1):

    # Setting training and validation Loss
    train_loss = []
    validation_loss = []
    tr_predictions = 0
    acc = 0
    val_predictions = 0

    #####
    # Train the model #
    #####
    model = model.train()
    for data in train_loader:

        # Moving tensors to GPU on CUDA enabled devices
        if device:
            input_ids, attention_mask, targets = data["input_ids"].cuda(), data["attention_mask"].cuda(), data["targets"]
        # Clear the gradients of variables
        optimizer.zero_grad()

    #### Forward pass
    # Pass input through the model
    output = model(
        input_ids=input_ids,
        attention_mask=attention_mask
    )
    # Compute batch loss
    loss = criterion(output, targets)
    # Convert output probabilities to class probabilities
    _, pred = torch.max(output, 1)
    # Track correct predictions
    tr_predictions += torch.sum(pred == targets)

    #### Backward Pass
    # Compute gradients wrt to model parameters
    loss.backward()
    # To avoid exploding gradients, we clip the gradients of the model
    nn.utils.clip_grad_norm_(model.parameters(), max_norm=1.0)
    # Perform parameter update
    optimizer.step()
    # Update Learning rate
    scheduler.step()
    # Update loss per mini batches
    train_loss.append(loss.item())

    #####
    # Validate the model #
    #####
    model.eval()
    with torch.no_grad():
        for data in valid_loader:

```

```

# Moving tensors to GPU on CUDA enabled devices
if device:
    input_ids, attention_mask, targets = data["input_ids"].cuda(), data["attention_mask"].cuda(), data["targets"].cuda()

#### Forward pass
# Pass input through the model
output = model(
    input_ids=input_ids,
    attention_mask=attention_mask
)
# Compute batch loss
loss = criterion(output, targets)
# Convert output probabilities to class probabilities
_, pred = torch.max(output, 1)
# Update loss per mini batches
validation_loss.append(loss.item())
# Track correct predictions
val_predictions += torch.sum(pred == targets)

```

```

# Compute accuracy
train_accuracy = tr_predictions.double()/len(train_sampler)
val_accuracy = val_predictions.double()/len(valid_sampler)

# Print loss statistics
print('Epoch: {}/{} \n\tTraining Loss: {:.6f} \n\tValidation Loss: {:.6f} \n\tTrain Accuracy: {:.6f} \n\tVal Accuracy: {:.6f}'.format(
    epoch, num_epochs, train_loss, validation_loss, train_accuracy, val_accuracy))

# Save model if validation loss is decreased
if val_accuracy > acc:
    print('Saving model...')
    torch.save(model.state_dict(), 'bert_base_fine_tuned.pt')
    acc = val_accuracy

```



```
Epoch: 1/10
  Training Loss: 0.478485
  Validation Loss: 0.426510
  Train Accuracy: 0.813221
  Val Accuracy: 0.848361
Saving model...
Epoch: 2/10
  Training Loss: 0.251598
  Validation Loss: 0.587404
  Train Accuracy: 0.912720
  Val Accuracy: 0.837432
Saving model...
Epoch: 3/10
  Training Loss: 0.147462
  Validation Loss: 0.694001
  Train Accuracy: 0.958333
  Val Accuracy: 0.848361
Saving model...
Epoch: 4/10
  Training Loss: 0.095958
  Validation Loss: 0.852052
  Train Accuracy: 0.976548
  Val Accuracy: 0.841530
Saving model...
Epoch: 5/10
  Training Loss: 0.062927
  Validation Loss: 0.967488
  Train Accuracy: 0.985504
  Val Accuracy: 0.842896
Saving model...
Epoch: 6/10
  Training Loss: 0.042360
  Validation Loss: 1.066000
  Train Accuracy: 0.990437
  Val Accuracy: 0.840164
Saving model...
Epoch: 7/10
  Training Loss: 0.032142
  Validation Loss: 1.132496
  Train Accuracy: 0.992410
  Val Accuracy: 0.833333
Saving model...
Epoch: 8/10
  Training Loss: 0.024429
  Validation Loss: 1.184951
  Train Accuracy: 0.993777
  Val Accuracy: 0.829235
Saving model...
Epoch: 9/10
  Training Loss: 0.018996
  Validation Loss: 1.230268
  Train Accuracy: 0.994991
  Val Accuracy: 0.831967
Saving model...
Epoch: 10/10
  Training Loss: 0.015075
  Validation Loss: 1.244014
  Train Accuracy: 0.995826
  Val Accuracy: 0.830601
Saving model...
```



```

# Track test loss
test_loss = 0.0
class_predictions = list(0. for i in range(3))
class_total = list(0. for i in range(3))
predictions = []
labels = []

model.eval()
with torch.no_grad():
    for data in test_loader:

        # Moving tensors to GPU on CUDA enabled devices
        if device:
            input_ids, attention_mask, targets = data["input_ids"].cuda(), data["attention_mask"].cuda(), data["targets"]

        ### Forward pass
        # Pass input through the model
        output = model(
            input_ids=input_ids,
            attention_mask=attention_mask
        )
        # Compute batch loss
        loss = criterion(output, targets)
        # Update loss
        test_loss += loss.item()
        # convert output probabilities to predicted class
        _, pred = torch.max(output, 1)

        predictions.extend(pred)
        labels.extend(targets)

predictions = torch.stack(predictions) if not device else torch.stack(predictions).cpu()
labels = torch.stack(labels) if not device else torch.stack(labels).cpu()

```

```

print(classification_report(predictions, labels, target_names=['neutral', 'positive', 'negative']))

```

	precision	recall	f1-score	support
neutral	0.68	0.71	0.69	146
positive	0.79	0.81	0.80	118
negative	0.92	0.90	0.91	468
accuracy			0.85	732
macro avg	0.79	0.81	0.80	732
weighted avg	0.85	0.85	0.85	732

```

cm = confusion_matrix(labels, predictions)
heatmap = sns.heatmap(cm, annot=True, fmt='d', cmap='Greens')
heatmap.yaxis.set_ticklabels(heatmap.yaxis.get_ticklabels(), rotation=0, ha='right')
heatmap.xaxis.set_ticklabels(heatmap.xaxis.get_ticklabels(), rotation=30, ha='right')
plt.xlabel('True sentiment')
plt.ylabel('Predicted sentiment');

```

